

Altube — Founder-Mode Beta Plan

Prepared June 2026. Lens: first-principles operator, pre-revenue beta, supply-first wedge (win AI creators), success = "creators upload here first."

Here's what I'd actually do. You will not out-engineer YouTube or Netflix on a beta budget. So don't try. Win on one thing: be **the platform AI creators trust and get paid better on**, with provenance and safety as the moat. Buy the hard ML, build the creator economics, and launch into ONE niche — not "the globe." Everything below is sequenced to prove that in ~6 months, not to build the final product.

0. The one decision that determines everything

You told me three things: pre-revenue beta, supply-first (win creators), and the beta must prove **creators choose you over YouTube**. That is the whole strategy. It means:

- **Don't optimize the viewer recommender first.** Viewers are downstream of content. If the best AI creators are here, viewers come. Spend your scarce energy on creator value, not on a Netflix-grade homepage no one has content to fill yet.
- **The "marketplace cold start"** is your #1 risk, not features. An empty platform with a beautiful UI is still empty. Solve supply with a hand-picked founding-creator cohort before you write a line of recommender code.
- **What most people miss:** a YouTube-killer doesn't win on the feed. It wins on a wedge YouTube structurally *can't* copy fast — here, that's **mandatory provenance + pre-publish safety + better creator monetization for synthetic media specifically**. YouTube has 2 billion users it can't risk alienating and a moderation system built for human-shot video. You can be AI-native from day one.

1. Positioning: the gap you can own *right now*

YouTube's blind spots with AI content, as of mid-2026:

1. **Trust is a mess.** AI slop, undisclosed deepfakes, and "is this real?" anxiety are rising. The EU AI Act Article 50 (enforcement begins **August 2026**) requires machine-readable

disclosure of AI-generated content. YouTube bolts labels on after upload. **You make verified provenance the entry ticket.**

2. **AI creators are second-class citizens** on YouTube — demonetization risk, unclear policy, no tools built for their workflow.
3. **Discovery rewards clickbait and watch-time**, not craft. Your multi-dimensional rating system (visual consistency, narrative, audio) can surface *quality* AI work that YouTube buries.

Your one-line position: *"The home for verified, high-craft AI video — where creators are paid fairly and viewers always know what they're watching."*

That sentence is your filter for every feature decision. If a feature doesn't serve "verified," "high-craft," "paid fairly," or "always know," it's post-beta.

2. Buy vs. build — the only architecture call that matters at beta

You cannot afford to build deepfake detection, an AI-or-not checker, content moderation, AND a recommender from scratch. Here's the honest split. **Buy everything that's a commodity; build only the creator-economics layer and your provenance UX, because that's your moat.**

Capability	Beta decision	Why
"Is this actually AI?" checker	Buy (Hive AI-Generated Content Classification, Sensity, or Sightengine)	These are mature commodity APIs returning confidence scores. Building your own is a research project.
Pre-publish safety (deepfake of real people, explicit, violence)	Buy (Hive Moderation 50+ categories, or AWS Rekognition <code>StartContentModeration</code>)	Same — solved problem. Wrap it in <i>your</i> policy logic.
Provenance / disclosure	Build the UX, adopt the standard (C2PA / Content Credentials)	This is your moat. The standard is free and adopted (Adobe, OpenAI, TikTok, camera makers). You build how it's surfaced and enforced.
Recommender	Buy/managed at beta (Amazon Personalize, Recombee, or Shaped), build later	A custom embedding recommender is 70k—400k+. At beta you don't have enough watch data to train one anyway. Managed gets you 80% there in days.
Video pipeline (transcode, store, deliver)	Buy (Mux, or Cloud Storage/S3 → Transcoder → CDN)	Mux is the fastest path; raw cloud stack is cheaper at scale. Start with Mux for the beta.
Multilingual	Buy (auto subtitle/translation APIs; UI i18n framework)	Don't build translation.
Creator monetization, ratings, credits, analytics	BUILD	This is the product. This is why creators choose you. Own it.

The principle: your engineering hours are worth most on the layer your competitor can't buy off a shelf. That's creator economics and provenance, not transcoding.

3. Content safety — your trust moat (do this right or nothing else matters)

One undisclosed deepfake of a real person going viral on your platform in week one kills the brand permanently. So safety is **gating, not garnish**.

Pre-publish pipeline (every upload, before it's ever visible):

1. **AI-or-not check** → confirm it's actually AI (your platform's whole premise). Human-shot video gets flagged/rerouted.
2. **Provenance check** → read C2PA / Content Credentials if present; record the generating model where disclosed.
3. **Safety scan** (Hive/Rekognition) → explicit, violence/gore, hate symbols, self-harm, weapons.
4. **Deepfake-of-real-person check** → the highest-risk category. Non-consensual likeness of a real, identifiable person = hard block, no appeal queue for obvious cases.
5. **Decision tree:**
 - High-confidence violation → **blocked before publish**, creator notified.
 - Ambiguous → **held in async human-review queue** (don't publish, don't auto-reject).
 - Clean → publish + provenance badge shown to viewers.

Trade-off to accept consciously: pre-publish blocking adds upload latency (minutes for video). At beta scale that's fine and the trust payoff is worth it. The hybrid model (instant publish + post-scan) is what big platforms do *because* they can't afford the latency at billions of uploads — you're not them, and "nothing bad ever published here" is a sharper brand than "we catch most of it eventually."

What most founders miss: the moderation *policy* and the *appeals UX* are harder than the ML. Write your policy doc before launch. Decide who reviews the ambiguous queue (you, at first — founders should do support and moderation personally in beta to learn the edge cases).

4. Product — what to build for the beta (and what to cut)

Build now (serves the wedge)

- **Frictionless verified upload** with the safety pipeline above + an automatic "**AI-Verified**" / **provenance badge**. This badge IS the product differentiator viewers see.
- **Multi-dimensional ratings** (Visual Consistency, Narrative/Script, Audio/Voice) → feeds a **Weighted Quality Score**. This is genuinely novel and serves "high-craft." Keep it.
- **Creator monetization the moment a creator qualifies**: subscriptions (creator sets price), tips. This is the "paid fairly" promise — make your revenue share visibly better than YouTube's and say so loudly.
- **Creator analytics dashboard** (views, watch time, rating breakdown, where viewers drop, geography). Creators live in analytics — this is a retention tool *for creators*.
- **Creator collaboration / joint-development** — a real wedge YouTube lacks. Even a lightweight "co-creator" credit + shared revenue split + a shared project space is a reason to come here. Prioritize a basic version.
- **AI help bot** for creators (upload help, policy questions, "why was I flagged"). Cheap to build on an LLM, high support-cost savings.

Defer (post-beta; build only after supply exists)

- The full **Netflix-style cinematic homepage** with hero spotlight + many algorithmic rows. Build a *simple* curated grid + the "Highly Rated" row at beta. The cinematic UI is a week-3 problem once you have content worth featuring.
- The **Altube Credits economy** (review rewards → unlock content / tip). Lovely idea, but it's a whole virtual-economy with abuse vectors (credit farming via fake reviews). Ship reviews *without* paid credits first; add the economy once you've seen real review behavior. **Risk most people miss**: paying for reviews above 150 chars trains people to write 151-character filler that passes the auto-check. Reward quality signals (helpful votes from creators), not length.
- **Community chat** between creators/viewers — valuable but a moderation surface you can't staff at beta. Phase it in creator-to-creator first (smaller, higher-trust group), then creator-to-viewer.
- **Paid subscription tiers unlocked at a creator threshold** (your attached-image idea) — good design, but it's a phase-2 monetization feature. Gate it behind proven creators so the first paid content is actually good.

The discipline here: every deferred item is deferred *because shipping it at beta either has no content to act on, or creates an abuse/moderation surface you can't yet staff*. That's the founder filter, not feature laziness.

5. The recommender — be honest about what's achievable

You asked for an algorithm "on par with / better than YouTube." Straight talk: **YouTube's recommender is the product of a decade and billions of dollars of watch data. You will not match it at beta, and you don't need to.** Here's the real path:

- **Beta:** managed recommendations (Amazon Personalize / Recombee / Shaped) + your **Weighted Quality Score** as a ranking signal. Your *advantage isn't algorithmic sophistication — it's better input signals*. YouTube ranks on watch-time (rewards clickbait). You rank partly on **multi-dimensional quality ratings**, which no one else has. That's how a simpler algorithm produces a *better feed* for craft content.
- **Dynamic dashboard** (genres shift based on recent watches): managed services support this out of the box via real-time events. Don't build it from scratch.
- **The throttle you described** (stop pushing content if engagement/ratings fall below a threshold): build this as a simple rules layer on top — it's a business rule, not ML. Easy win, do it.
- **Phase 2 (post product-market-fit, with real data):** build custom embeddings (video + transcript via models like Gemini Embedding / Cohere / open-weight Qwen3) and vector similarity. Only when recommendations are a proven revenue pillar and you have the data to train on.

What "better than YouTube" actually means for you: not a smarter black box — a feed that surfaces quality and respects disclosure, plus a creator-side promise that good work won't get buried by clickbait. That's a positioning win, achievable now.

6. Multilingual & global — design for it, don't boil the ocean

- **Internationalize from day one** (i18n framework, externalized strings, RTL-ready, Unicode everywhere). Retrofitting this later is painful — cheap to do now.
- **Auto-generate subtitles + translations** per video via speech-to-text + translation APIs. AI video is uniquely well-suited to this — many have synthetic voiceover that transcribes

cleanly.

- **But launch the *audience* in one or two languages/regions first.** Supporting the infrastructure globally $\neq$ marketing globally. Pick your beta niche's primary language. Global-ready architecture, focused go-to-market.

7. Architecture — scalable, web-first, app-ready

Principle: web-first, but build the API so a mobile app is a client, not a rewrite.

CLIENTS

Web: Next.js (SSR for SEO on public video pages – critical for discovery)

Mobile (Phase 2): React Native or Flutter – reuses the same API

API LAYER

Single API gateway → serverless functions (start) → containers (scale)

GraphQL or REST; versioned. This is the contract every client uses.

CORE SERVICES (serverless to start, extract to services as they grow)

- Auth & profiles
- Upload orchestration → triggers Safety Pipeline
- Safety Pipeline (async): AI-check → provenance → moderation → deepfake → publish/hold/block
- Ratings & Weighted Quality Score
- Monetization (subscriptions, tips, payouts) [Stripe/Stripe Connect]
- Recommendations (managed service adapter)
- Analytics ingestion
- Notifications / AI help bot

DATA

- Primary metadata/users/reviews: NoSQL (Firestore / MongoDB) – flexible, fast to iterate
- Money/transactions: relational (Postgres) – you want ACID + auditability for payments.

Don't put payouts in NoSQL.

- Search: managed search index (e.g., Algolia / OpenSearch)
- Events/analytics: event stream → warehouse (for the recommender + creator analytics later)

VIDEO PIPELINE

Upload → object storage (S3 / GCS) → Mux (beta) or Transcoder API → CDN → adaptive streaming (HLS)

Safety scan runs on the stored asset BEFORE the CDN/publish step.

3RD-PARTY (buy)

- Hive / Sensity / Sightengine (AI + deepfake + moderation)
- C2PA / Content Credentials SDK (provenance)
- Amazon Personalize / Recombee / Shaped (recommendations)
- Stripe Connect (creator payouts)
- Speech-to-text + translation (multilingual)

Scalability decisions that matter and why:

- **Serverless first** = no ops overhead, scales with spiky video traffic, you pay for what you use. Extract hot paths to containers later. Don't pre-optimize for scale you don't have.
- **Async safety pipeline via a queue** = uploads don't block the user-facing app; you can add review steps without re-architecting.

- **Split datastores by job** (NoSQL for content, SQL for money). The single most common architecture mistake here is putting payments in a document DB for "consistency" of stack. Don't.
 - **API-as-contract** = the mobile app, future TV apps, and partners all consume the same layer. This is what makes you "available as an app" without a rewrite.
-

8. Sequenced roadmap (the part most plans skip)

Phase 0 — Pre-build (Weeks 0–4): solve supply before you build the feed

- Hand-recruit **15–30 founding AI creators** (the kind who already post AI film/anime/sci-fi). Offer: better revenue share, the verified badge, founding-creator status, direct line to you. This cohort IS your beta.
- Write the **content policy + moderation playbook**. Decide your hard-block vs. review-queue lines.
- Lock buy-vs-build vendor choices; get API keys; test the safety pipeline on real AI clips.

Phase 1 — Closed beta (Weeks 4–12): prove creators upload here first

- Ship: verified upload + safety pipeline + verified badge, profiles, simple curated grid + "Highly Rated" row, multi-dimensional ratings, creator analytics dashboard, AI help bot.
- Monetization MVP: subscriptions + tips via Stripe Connect for founding creators.
- **The metric that matters:** do founding creators post *new/exclusive* work here, and come back to check analytics? That's your "choose us over YouTube" proof. Talk to all 30 of them weekly.

Phase 2 — Open beta (Weeks 12–24): bring viewers

- Managed recommender + dynamic dashboard + the engagement throttle.
- Netflix-style homepage (now you have content to feature).
- Multilingual subtitles/translations; open to a second language/region.
- Creator collaboration / joint-dev v1.
- Begin the investor conversation here, with real creator-retention and early viewer-retention data — *that's* the pitch, not the feature list.

Phase 3 — Post-PMF (Month 6+): the things you deferred

- Credits economy (with abuse controls), community chat (phased), custom embedding recommender, mobile apps, paid creator-threshold tiers, more regions.

9. What breaks first, and the risks most people miss

1. **Cold start (highest risk).** No content → no viewers → no content. *Mitigation:* Phase 0 founding-creator cohort before any feed work. If you can't get 15 good creators to commit, the product thesis is wrong — find that out in week 2, cheaply.
 2. **One bad deepfake destroys trust permanently.** *Mitigation:* pre-publish hard-block for non-consensual real-person likeness; founders personally review the ambiguous queue at beta.
 3. **The credits economy gets gamed** (fake reviews for credits). *Mitigation:* don't ship it at beta; reward quality signals, not character count.
 4. **"Better than YouTube's algorithm" is the wrong goal** and will sink your time. *Mitigation:* compete on *signal quality* (your ratings) and *positioning*, not algorithmic horsepower. Managed recommender until you have data.
 5. **Building four hard ML systems at once with a beta team.** *Mitigation:* buy all of them. Build only creator economics + provenance UX.
 6. **Moderation cost/latency at scale.** Fine at beta, plan for it: budget the per-minute API cost now so it doesn't surprise you at 10x volume.
 7. **Regulatory tailwind is also a requirement.** EU AI Act Article 50 disclosure (Aug 2026) — your provenance-first design turns a compliance burden into your differentiator. Don't treat it as optional.
-

10. What the top 3% do differently

- They **launch embarrassingly narrow** (one niche, 30 creators) and dominate it, instead of building "global multilingual everything" first.
- They **do moderation and creator support by hand** in beta — founders in the trenches — to learn the edge cases code can't anticipate yet.
- They **buy boring infrastructure** and pour every original engineering hour into the one or two things that are actually their moat. For you: creator economics + verified provenance.

- They **pick a sharp enemy and a sharp promise**. "Verified, high-craft AI video where creators are paid fairly and viewers always know what they're watching." Everything serves that sentence.

Sources

- [Hive AI-Generated & Deepfake Content Classification](https://thehive.ai/apis/ai-generated-content-classification) (https://thehive.ai/apis/ai-generated-content-classification)
- [Sensity AI — Deepfake Detection](https://sensity.ai/) (https://sensity.ai/)
- [Sightengine Deepfake Detection API](https://sightengine.com/docs/deepfake-detection) (https://sightengine.com/docs/deepfake-detection)
- [AWS Rekognition Content Moderation](https://aws.amazon.com/rekognition/content-moderation/) (https://aws.amazon.com/rekognition/content-moderation/)
- [Eden AI — Top Explicit Content Detection APIs 2026](https://www.edenai.co/post/top-10-explicit-content-detection-apis) (https://www.edenai.co/post/top-10-explicit-content-detection-apis)
- [Amazon Personalize](https://aws.amazon.com/personalize/) (https://aws.amazon.com/personalize/)
- [Recombee — Real-time Recommendation Engine](https://www.recombee.com/product) (https://www.recombee.com/product)
- [Shaped — Recommendation APIs](https://www.shaped.ai/blog/5-best-apis-for-adding-personalized-recommendations-to-your-app-in-2025) (https://www.shaped.ai/blog/5-best-apis-for-adding-personalized-recommendations-to-your-app-in-2025)
- [Mux — Building a video recommendation engine with AI embeddings](https://www.mux.com/docs/examples/ai-recommendation-engine) (https://www.mux.com/docs/examples/ai-recommendation-engine)
- [Content Authenticity Initiative — State of Content Authenticity 2026](https://contentauthenticity.org/blog/the-state-of-content-authenticity-in-2026) (https://contentauthenticity.org/blog/the-state-of-content-authenticity-in-2026)
- [C2PA Adoption Status 2026 \(EU AI Act Article 50\)](https://www.eyesift.com/faq/c2pa-content-credentials-2026-cryptographic-provenance-adoption/) (https://www.eyesift.com/faq/c2pa-content-credentials-2026-cryptographic-provenance-adoption/)